







http://aima.cs.berkeley.edu/)

DSE CL 557 - Artificial and Computational Intelligence

**#3. Problem Solving Agent using Uninformed Search** 

Sunday, September 13, 2020

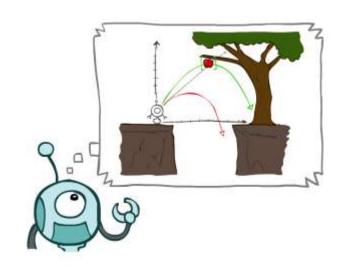


Dr. Saikishor Jangiti

### Today



- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
  - Depth-First Search
  - Breadth-First Search
  - Uniform-Cost Search
  - Depth Limited Search,
  - Iterative Deepening Search,
  - Bidirectional Search

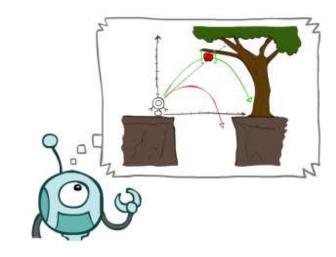


Reference: Chapter 3.1-3.4 from AI: A modern approach (Russell, Norvig)

### Agents that Plan



- How an agent can act by establishing *goals*
- Considering sequences of actions that might achieve those goals.

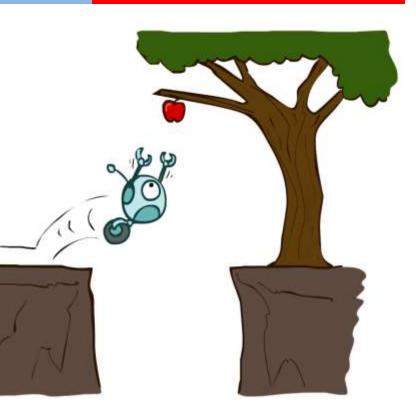


• A goal and a set of means for achieving the goal is called a *problem*, and the process of exploring what the means can do is called *search*.

### Reflex Agents

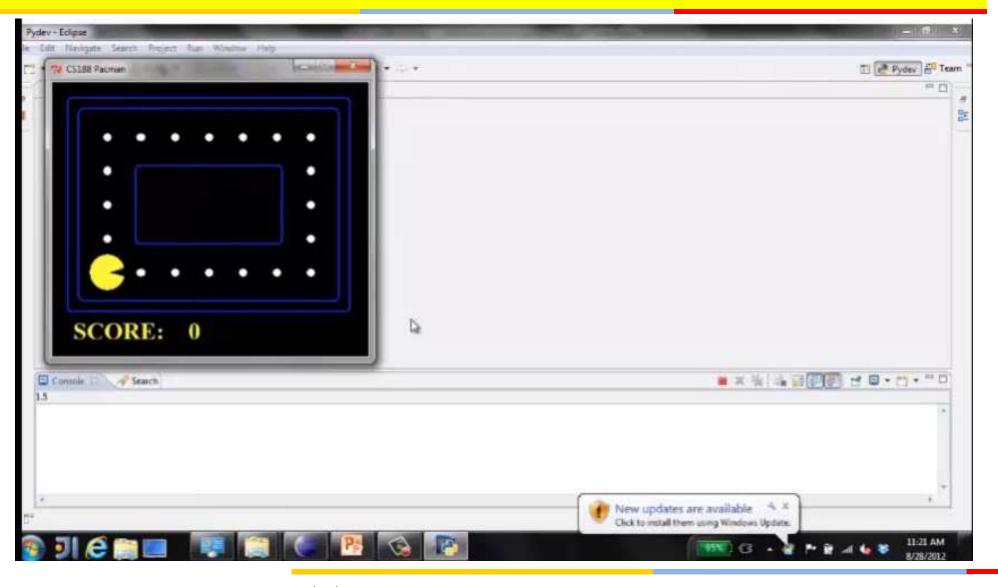


- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world is
- Can a reflex agent be rational?



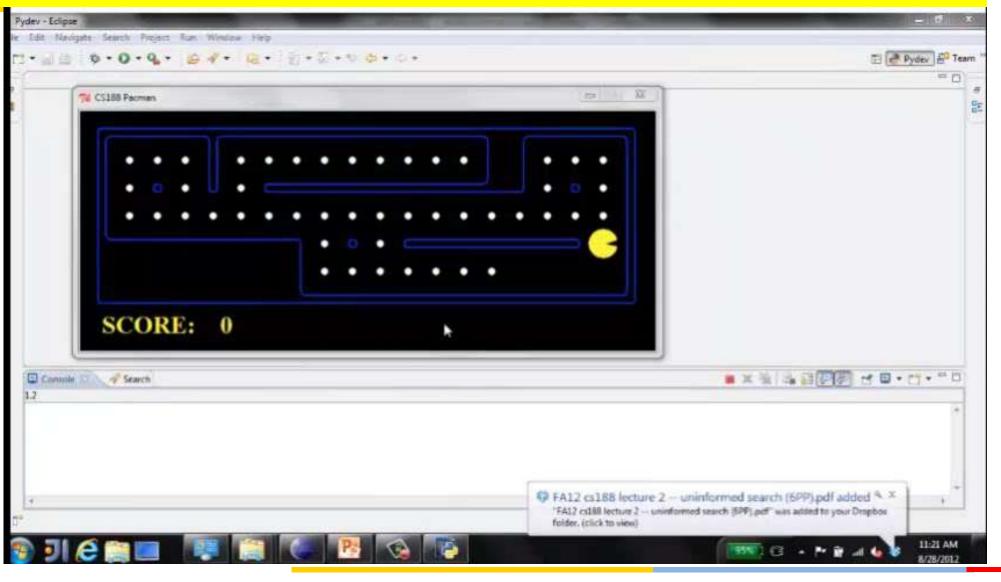
## Video of Demo Reflex Optimal





#### Video of Demo Reflex Odd

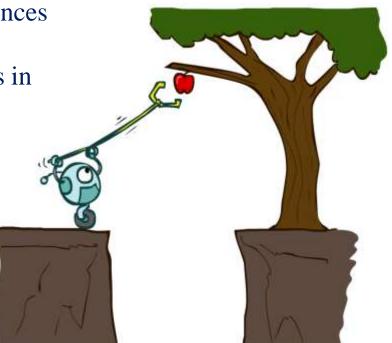




### Planning Agents

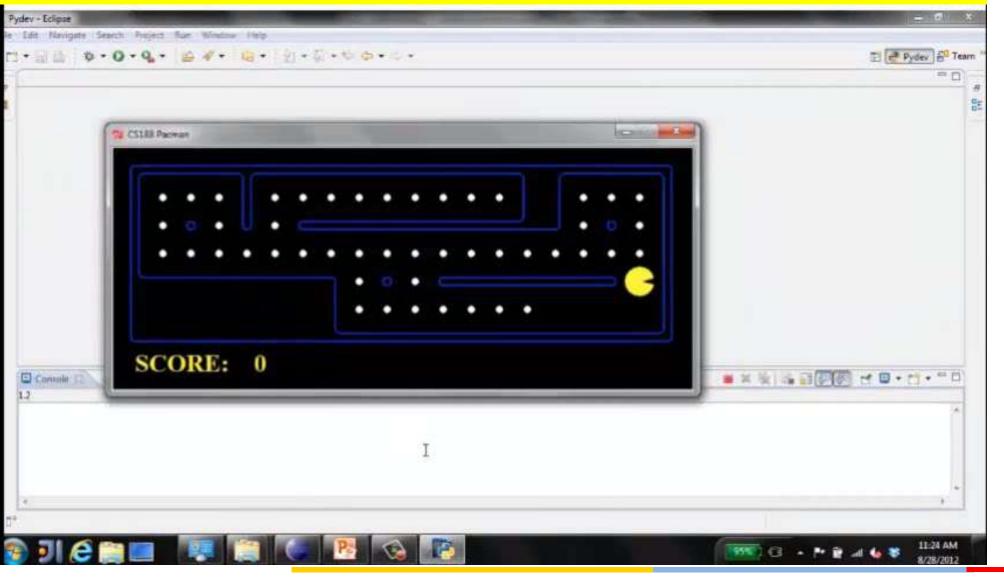


- Planning agents:
  - Ask "what if"
  - Decisions based on (hypothesized) consequences of actions
  - Must have a model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - Consider how the world WOULD BE
- Optimal vs. complete planning
- Planning vs. replanning



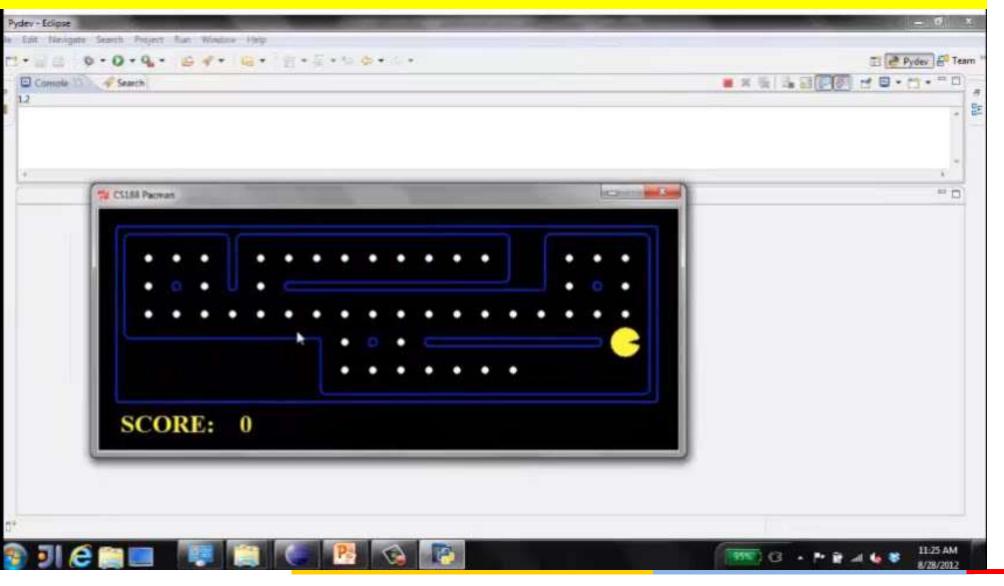
## Video of Demo Replanning





### Video of Demo Mastermind





### Search Problems





#### Search Problems



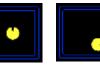
- A search problem consists of:
  - A state space





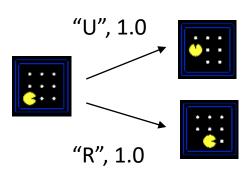








 A successor function (with actions, costs)



- A start state and a goal test
- A solution is a sequence of actions (a plan) which transforms the start state to a goal state

11

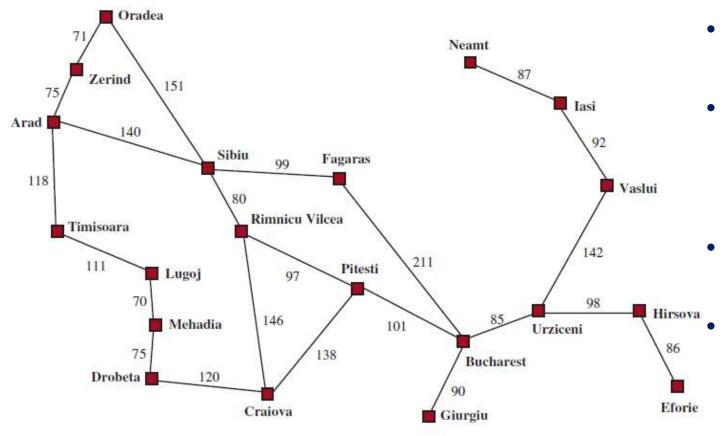
### Search Problems Are Models





### Example: Traveling in Romania



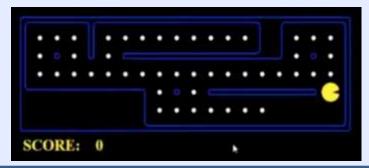


- State space:
  - Cities
- Successor function:
  - Roads: Go toadjacent city withcost = distance
- Start state:
  - Arad
  - Goal test:
    - Is state ==
       Bucharest?
- Solution?

## What's in a State Space?



The world state includes every last detail of the environment



A search state keeps only the details needed for planning (abstraction)

- Problem: Pathing
  - States: (x,y) location
  - Actions: UDLR
  - Successor: update location only
  - Goal test: is (x,y)=END

- Problem: Eat-All-Dots
  - States:  $\{(x,y), dot booleans\}$
  - Actions: UDLR
  - Successor: update location and possibly a dot boolean
  - Goal test: dots all false

## State Space Sizes?



#### • World state:

- Agent positions: 120

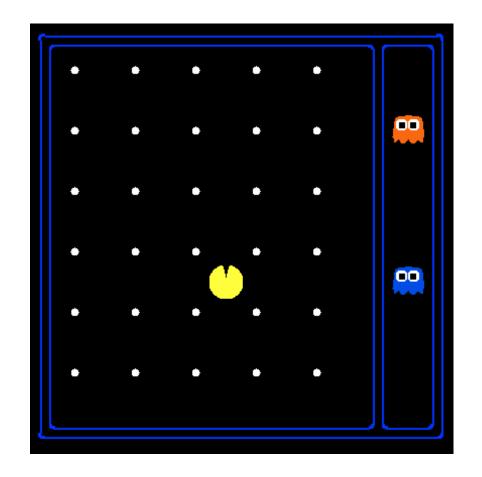
- Food count: 30

- Ghost positions: 12

Agent facing: UDLR

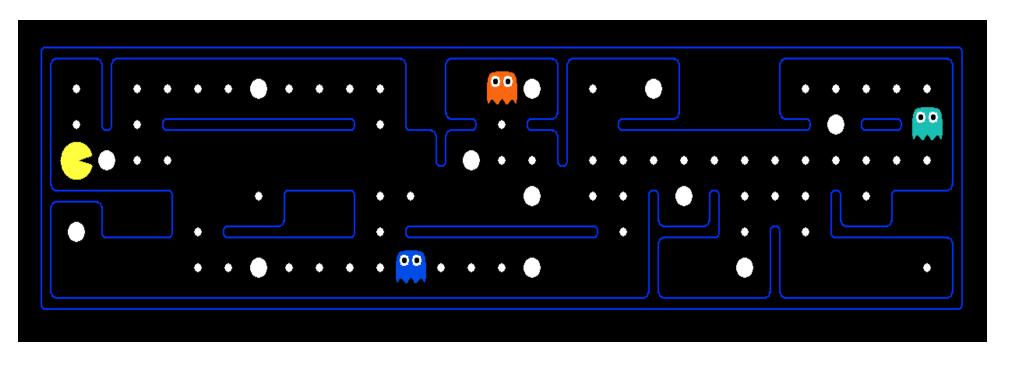
#### How many

- World states?
   120x(2<sup>30</sup>)x(12<sup>2</sup>)x4
- States for pathing?120
- States for eat-all-dots?120x(2<sup>30</sup>)



### Quiz: Safe Passage

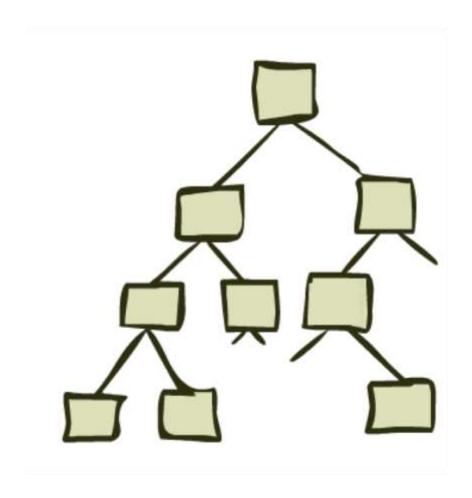




- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify?
  - (agent position, dot booleans, power pellet booleans, remaining scared time)

## State Space Graphs and Search Trees

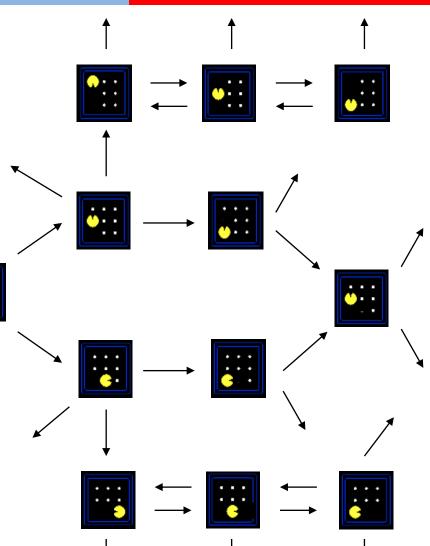




### State Space Graphs

BITS Pilani
Pilani Campus

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



### State Space Graphs

b

S

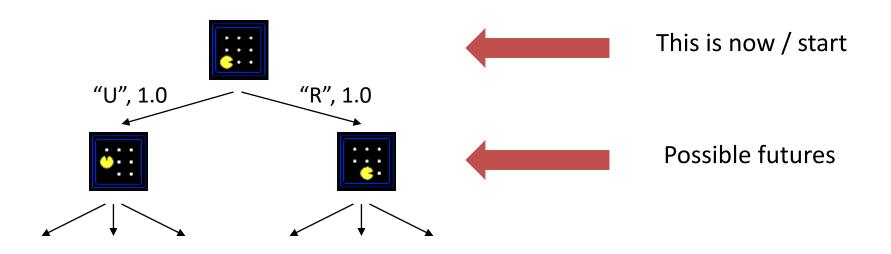


- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea

Tiny state space graph for a tiny search problem

#### Search Trees





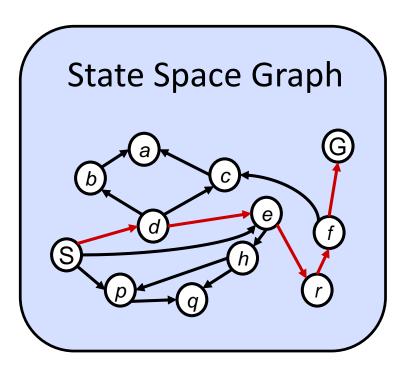
#### • A search tree:

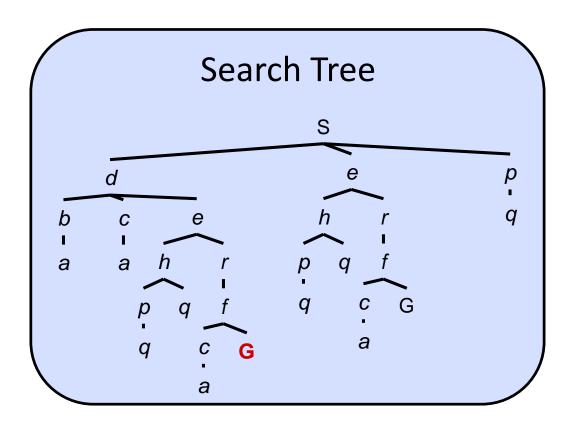
- A "what if" tree of plans and their outcomes
- The start state is the root node
- Children correspond to successors
- Nodes show states, but correspond to PLANS that achieve those states
- For most problems, we can never actually build the whole tree

## State Space Graphs vs. Search Trees



Each NODE in the search tree is an entire PATH in the state space graph.





We construct both on demand –

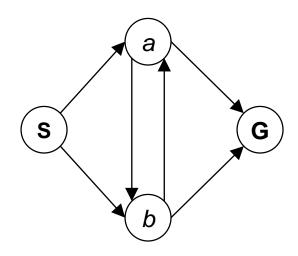
and we construct as little as possible.

### Quiz: State Space Graphs vs. Search Trees



Consider this 4-state graph:

How big is its search tree (from S)?






22

### Quiz: State Space Graphs vs. Search Trees



Consider this 4-state graph:

How big is its search tree (from S)?




Important: Lots of repeated structure in the search tree!

23

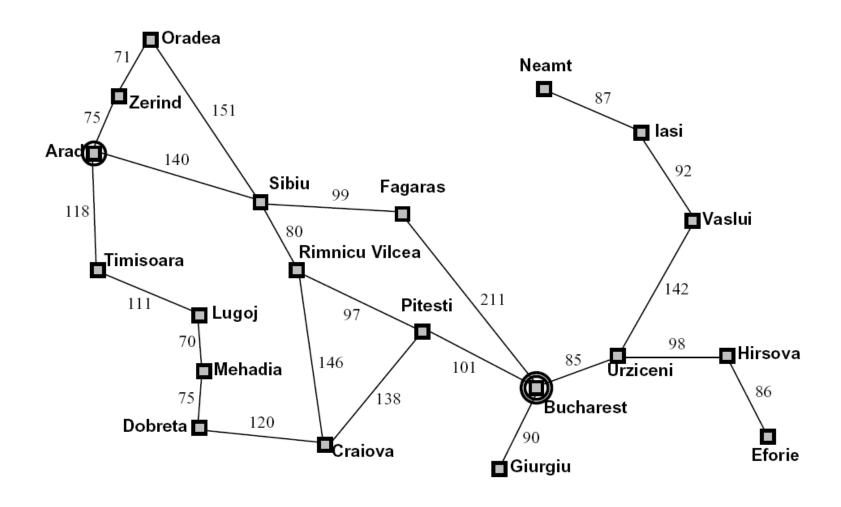
### Tree Search





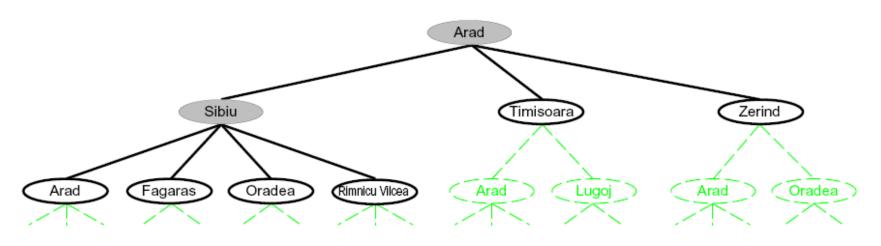
### Search Example: Romania





### Searching with a Search Tree





- Search:
  - -Expand out potential plans (tree nodes)
  - -Maintain a fringe of partial plans under consideration
  - -Try to expand as few tree nodes as possible

26

#### General Tree Search



```
function TREE-SEARCH( problem, strategy) returns a solution, or failure initialize the search tree using the initial state of problem loop do

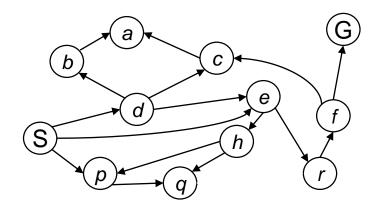
if there are no candidates for expansion then return failure choose a leaf node for expansion according to strategy

if the node contains a goal state then return the corresponding solution else expand the node and add the resulting nodes to the search tree end
```

- Important ideas:
  - Fringe
  - Expansion
  - Exploration strategy
- Main question: which fringe nodes to explore?

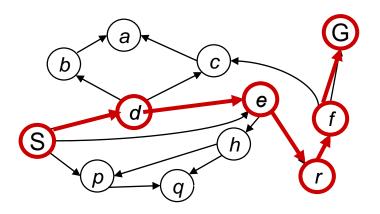
# Example: Tree Search

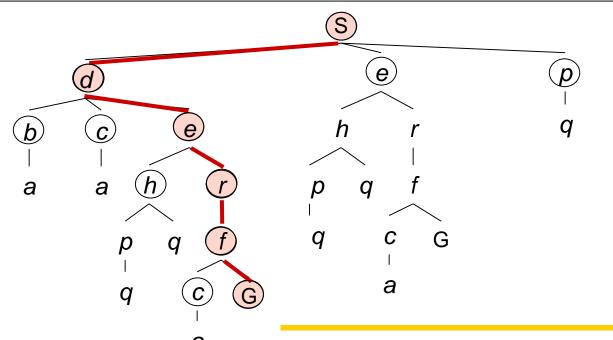


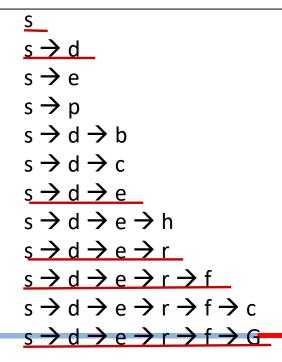


### Example: Tree Search



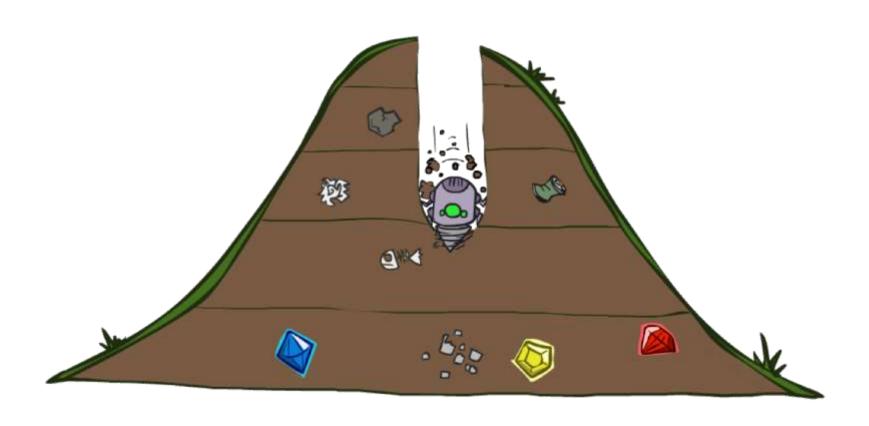






# Depth-First Search



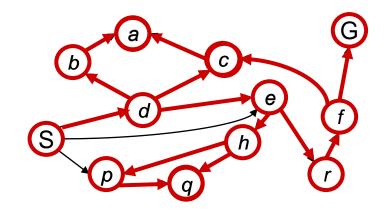


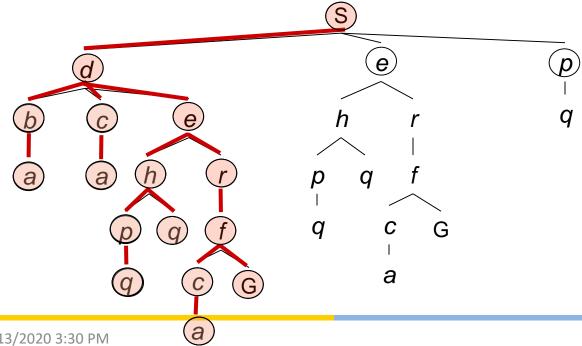
### Depth-First Search



Strategy: expand a deepest node first

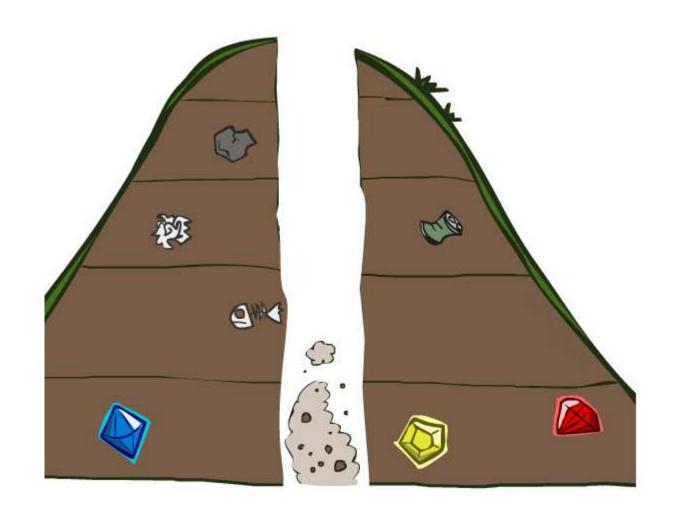
*Implementation:* Fringe is a LIFO stack





## Search Algorithm Properties

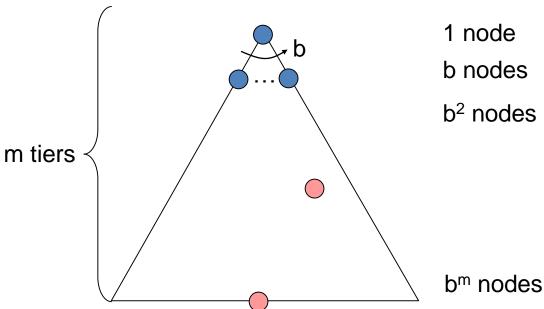




## Search Algorithm Properties



- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
  - b is the branching factor
  - m is the maximum depth
  - solutions at various depths



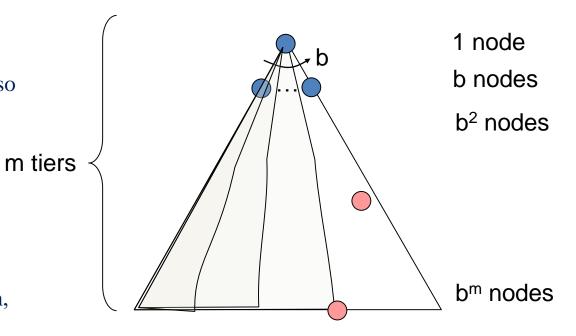
• Number of nodes in entire tree?

$$-1+b+b^2+....b^m = O(b^m)$$

### Depth-First Search (DFS) Properties

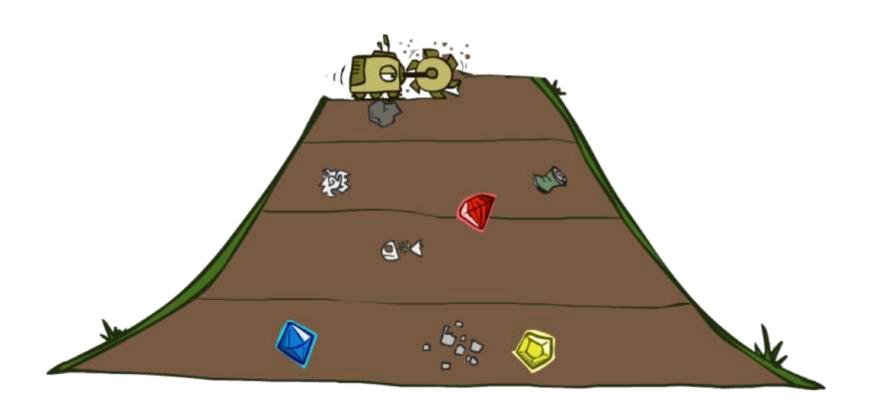


- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If m is finite, takes time O(b<sup>m</sup>)
- How much space does the fringe take?
  - Only has siblings on path to root, so O(bm)
- Is it complete?
  - m could be infinite, so only if we prevent cycles (more later)
- Is it optimal?
  - No, it finds the "leftmost" solution, regardless of depth or cost



### Breadth-First Search





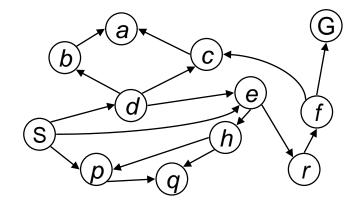
#### Breadth-First Search

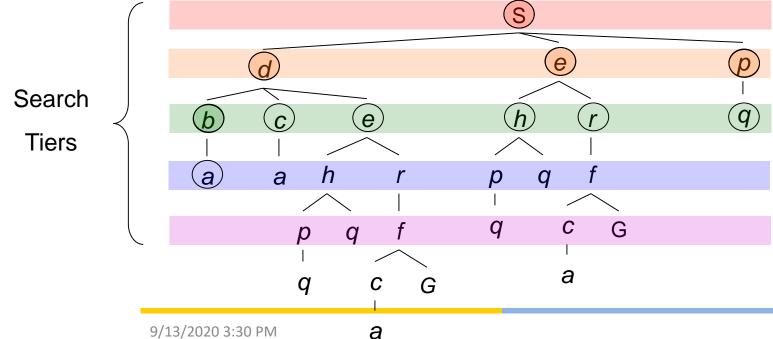


Strategy: expand a shallowest node first

*Implementation: Fringe* 

is a FIFO queue





## Breadth-First Search (BFS) Properties



- What nodes does BFS expand?
  - Processes all nodes above shallowest solution
  - Let depth of shallowest solution be s
  - Search takes time O(b<sup>s</sup>)

How much space does the fringe take?

Has roughly the last tier, so O(b<sup>s</sup>)

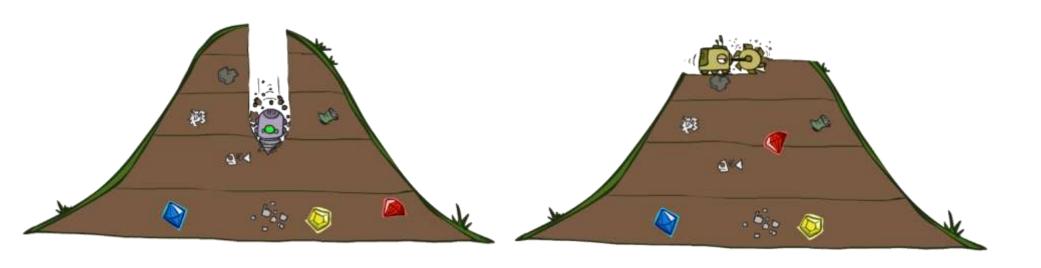
- Is it complete?
  - s must be finite if a solution exists, so yes!
- Is it optimal?
  - Only if costs are all 1 (more on costs later)

b nodes
b² nodes
b² nodes
bs nodes
bm nodes

1 node

# Quiz: DFS vs BFS





# Video of Demo Maze Water DFS/BFS (part 1)





# Video of Demo Maze Water DFS/BFS (part 2)





#### Quiz: DFS vs BFS



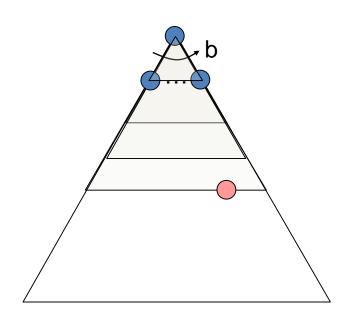
• When will BFS outperform DFS?

• When will DFS outperform BFS?

# Iterative Deepening

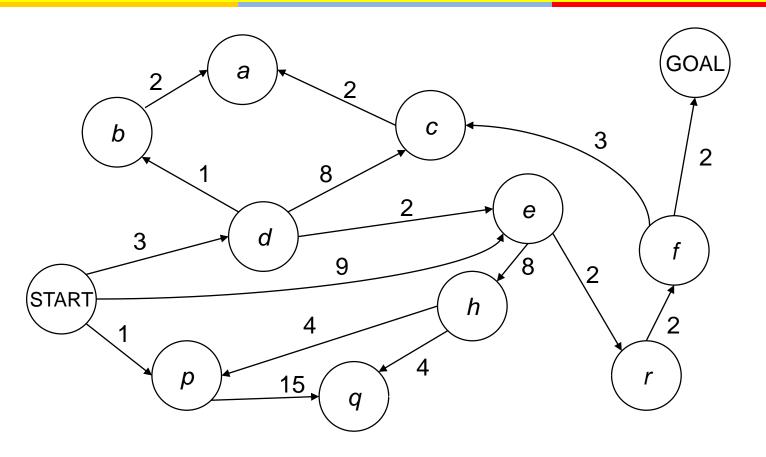


- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution...
  - Run a DFS with depth limit 3. .....
- Isn't that wastefully redundant?
  - Generally most work happens in the lowest level searched, so not so bad!



#### **Cost-Sensitive Search**





BFS finds the shortest path in terms of number of actions. It does not find the least-cost path. We will now cover a similar algorithm which does find the least-cost path.

# **Uniform Cost Search**



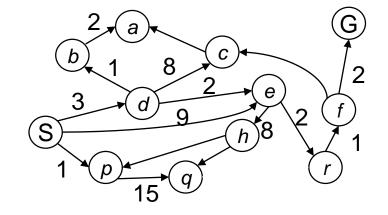


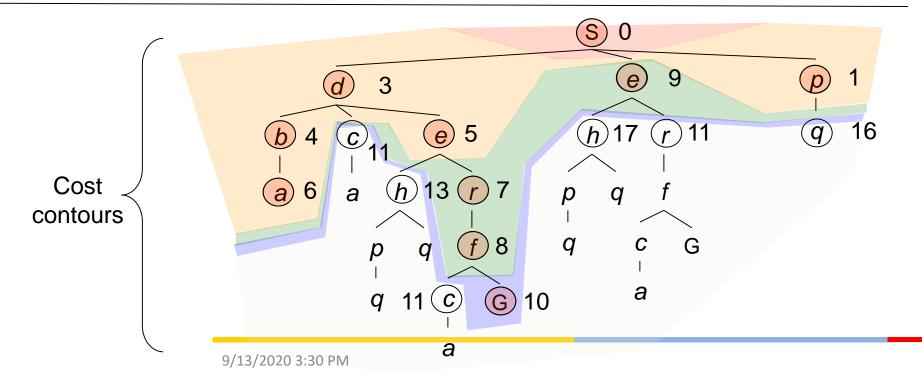
#### **Uniform Cost Search**



Strategy: expand a cheapest node first:

Fringe is a priority queue (priority: cumulative cost)

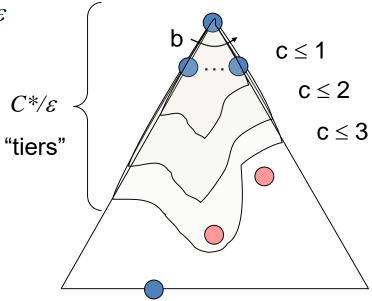




# Uniform Cost Search (UCS) Properties



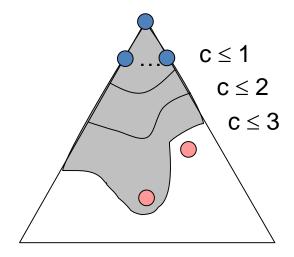
- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\varepsilon$ , then the "effective depth" is roughly  $C^*/\varepsilon$
  - Takes time  $O(b^{C^*/\varepsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C*/\varepsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (Proof next lecture via A\*)



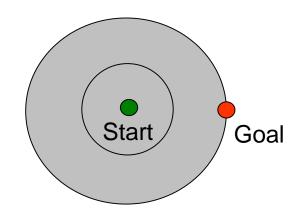
#### **Uniform Cost Issues**



- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!



- The bad:
  - Explores options in every "direction"
  - No information about goal location
- We'll fix that soon!



# Video of Demo Empty UCS





# Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 1)





# Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 2)





#### Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 3)





51

#### The One Queue



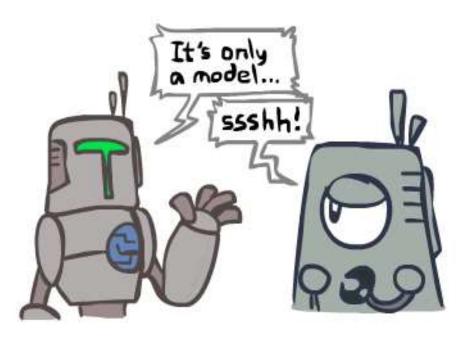
- All these search algorithms are the same except for fringe strategies
  - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
  - Practically, for DFS and BFS, you ca avoid the log(n) overhead from an actual priority queue, by using stacks and queues
  - Can even code one implementation that takes a variable queuing object



#### Search and Models

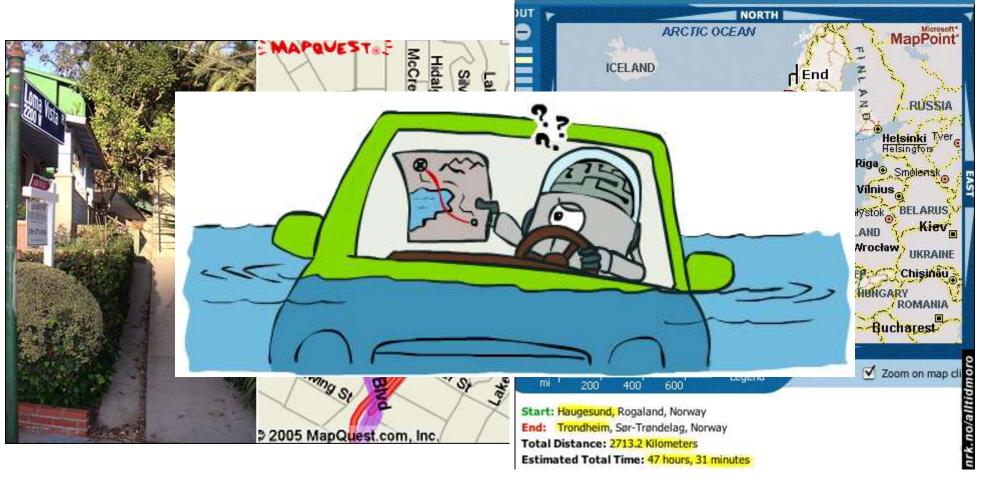


- Search operates over models of the world
  - The agent doesn't actually try all the plans out in the real world!
  - Planning is all "in simulation"
  - Your search is only as good as your models...



## Search Gone Wrong?





#### Thank You





# Any more Queries