





#### BITS Pilani Pilani Campus



#### DSE CL 557 - Artificial and Computational Intelligence

**#3. Informed Search** 

Sunday, September 20, 2020



Dr. Saikishor Jangiti

# Today



• Heuristics

Greedy Search

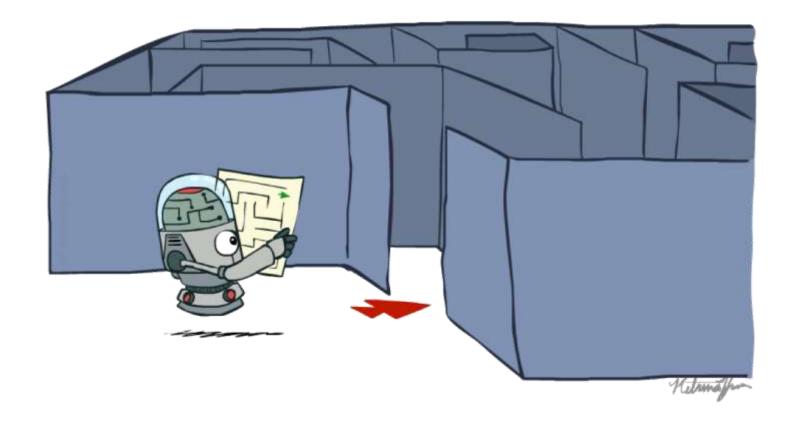
• A\* Search



Reference: Chapter 3.5 & 3.6 from AI: A modern approach (Russell, Norvig)

# Recap: Search





## Recap: Search



#### • Search problem:

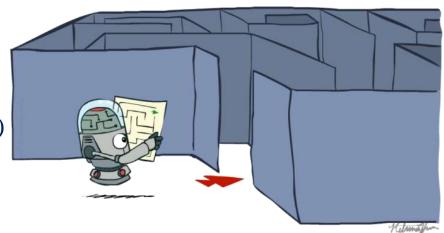
- States (configurations of the world)
- Actions and costs
- Successor function (world dynamics)
- Start state and goal test



- Nodes: represent plans for reaching states
- Plans have costs (sum of action costs)

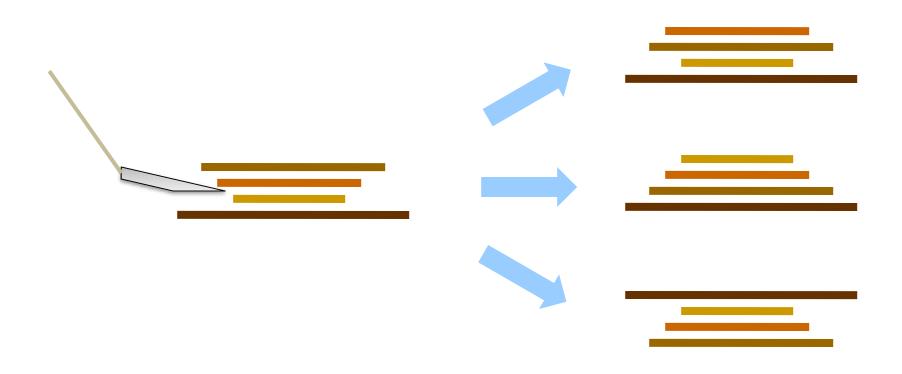
#### • Search algorithm:

- Systematically builds a search tree
- Chooses an ordering of the fringe (unexplored nodes)
- Optimal: finds least-cost plans



# Example: Pancake Problem





Cost: Number of pancakes flipped

## Example: Pancake Problem



#### BOUNDS FOR SORTING BY PREFIX REVERSAL

William H. GATES

Microsoft, Albuquerque, New Mexico

Christos H. PAPADIMITRIOU\*†

Department of Electrical Engineering, University of California, Berkeley, CA 94720, U.S.A.

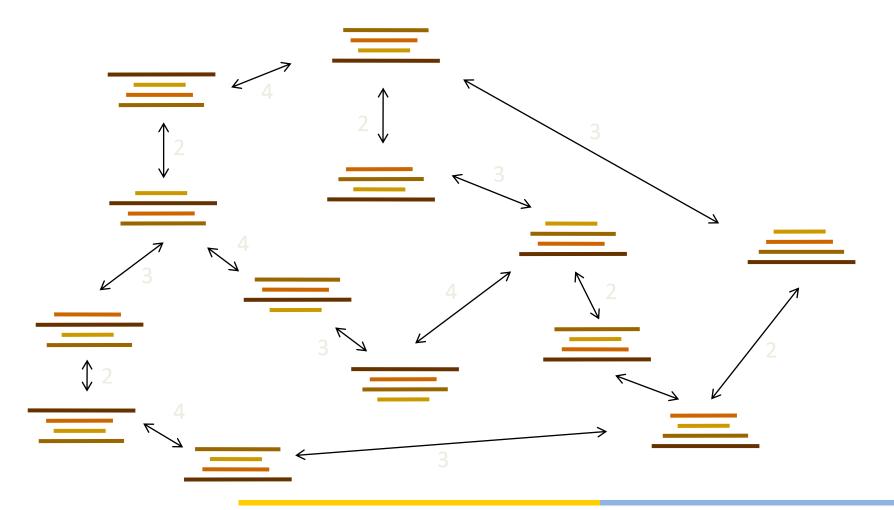
Received 18 January 1978 Revised 28 August 1978

For a permutation  $\sigma$  of the integers from 1 to n, let  $f(\sigma)$  be the smallest number of prefix reversals that will transform  $\sigma$  to the identity permutation, and let f(n) be the largest such  $f(\sigma)$  for all  $\sigma$  in (the symmetric group)  $S_n$ . We show that  $f(n) \leq (5n+5)/3$ , and that  $f(n) \geq 17n/16$  for n a multiple of 16. If, furthermore, each integer is required to participate in an even number of reversed prefixes, the corresponding function g(n) is shown to obey  $3n/2 - 1 \leq g(n) \leq 2n + 3$ .

# Example: Pancake Problem



### State space graph with costs as weights



#### **General Tree Search**



function TREE-SEARCH (problem, strategy) returns a solution, or failure initialize the search tree using the initial state of problem loop do if there are no candidates for expansion then return failure choose a leaf node for expansion according to strategy if the node contains a goal state then return the corresponding solution else expand the node and add the resulting nodes to the search tree end Action: flip top two Path to reach goal: Cost: 2 Flip four, flip three Total cost: 7

## The One Queue



- All these search algorithms are the same except for fringe strategies
  - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
  - Practically, for DFS and BFS, you can avoid the log(n) overhead from an actual priority queue, by using stacks and queues
  - Can even code one implementation that takes a variable queuing object



# **Uninformed Search**



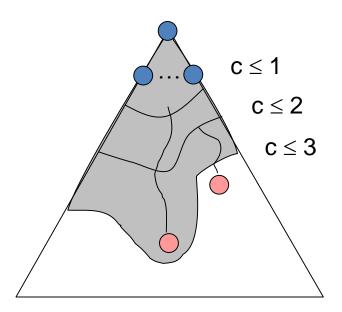


## **Uniform Cost Search**

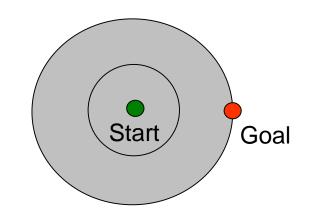


• Strategy: expand lowest path cost

• The good: UCS is complete and optimal!



- The bad:
  - Explores options in every "direction"
  - No information about goal location



# Video of Demo Contours UCS Empty





# Video of Demo Contours UCS Pacman Small Maze





# **Informed Search**



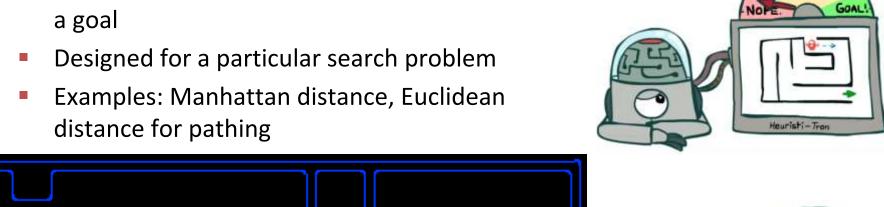



## Search Heuristics




#### A heuristic is:

A function that *estimates* how close a state is to a goal

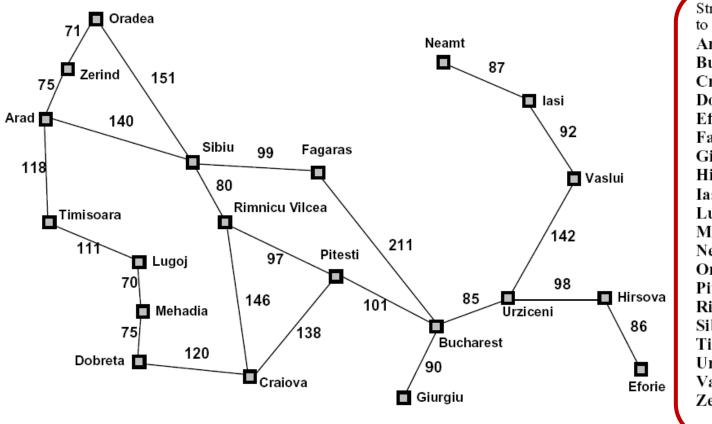






# **Example: Heuristic Function**






Straight-line distar to Bucharest	ice
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

h(x)

# **Example: Heuristic Function**




Heuristic: the number of the largest pancake that is still out of place



# **Greedy Search**






# **Example: Heuristic Function**




lead



Straight-line distance to Bucharest Arad 366 Bucharest Craiova 160 Dobreta 242 Eforie 161 Fagaras 178 Giurgiu 77 Hirsova 151 226 Lugoj 244 Mehadia 241 Neamt 234 Oradea 380 Pitesti 98 Rimnicu Vilcea 193 Sibiu 253 Timisoara 329 Urziceni 80 Vaslui 199 Zerind 374

h(x)

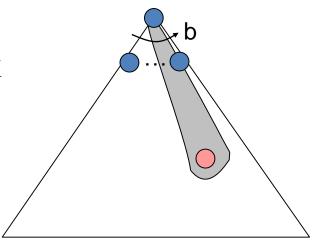




## **Greedy Search**

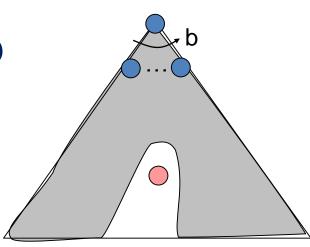
BITS Pilani
Pilani Campus

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state



- A common case:
  - Best-first takes you straight to the (wrong) goal

Worst-case: like a badly-guided DFS



# Video of Demo Contours Greedy (Empty)





# Video of Demo Contours Greedy (Pacman Small Maze)





## A\* Search





# A\* Search



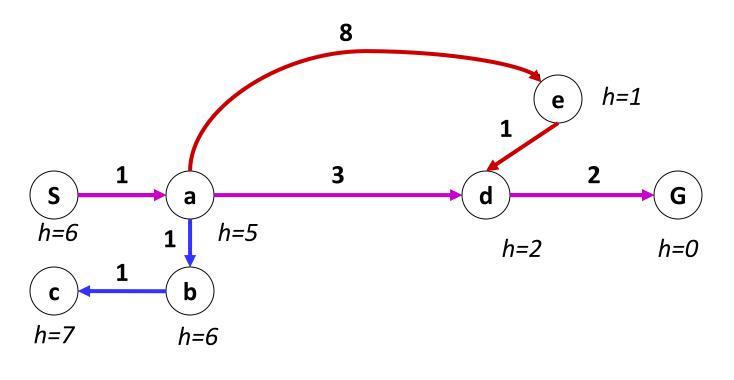


U

## Combining UCS and Greedy



- Uniform-cost orders by path cost, or backward cost g(n)
- Greedy orders by goal proximity, or forward cost h(n)



g = 0h=6 g = 1h=5 g = 9g = 2g = 4d b h=1 h=6 h=2 g = 3*g* = 6 g = 10h=7h=0h=2 g = 12h=0


A\* Search orders by the sum: f(n) = g(n) + h(n)

Example: Teg Grenager

### When should A\* terminate?



Should we stop when we enqueue a goal?




No: only stop when we dequeue a goal

27

# Is A\* Optimal?






- What went wrong?
- Actual bad goal cost < estimated good goal cost</li>
- We need estimates to be less than actual costs!

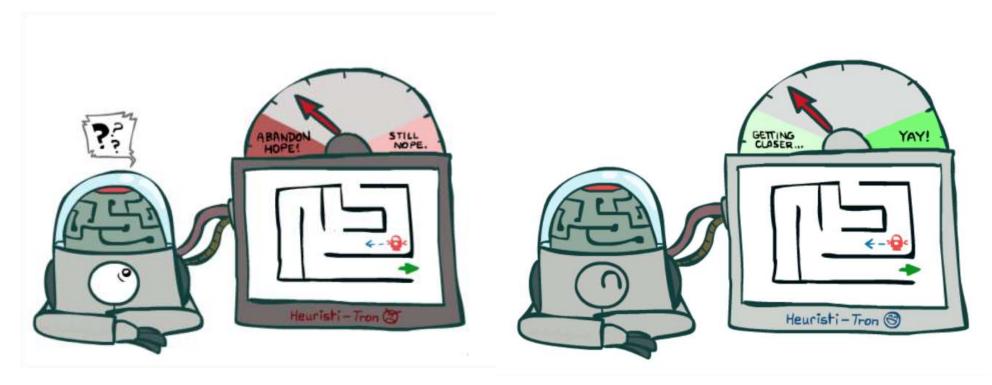
## Admissible Heuristics





# Idea: Admissibility





Inadmissible (pessimistic)
heuristics break optimality by
trapping good plans on the fringe

Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

## Admissible Heuristics



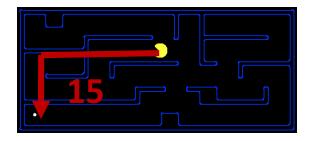
31

• A heuristic *h* is *admissible* (optimistic) if:

$$0 \le h(n) \le h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

• Examples:

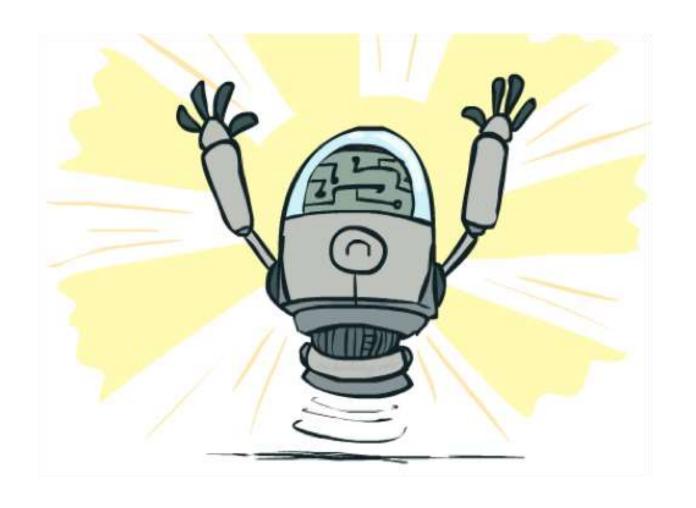




• Coming up with admissible heuristics is most of what's involved in using A\* in practice.

# Optimality of A\* Tree Search





## Optimality of A\* Tree Search

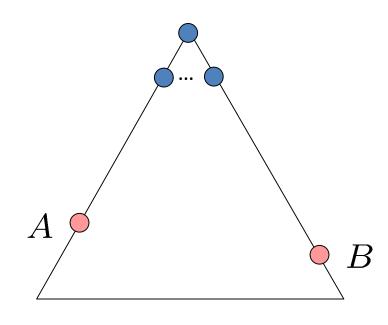


#### Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

#### Claim:

• A will exit the fringe before B

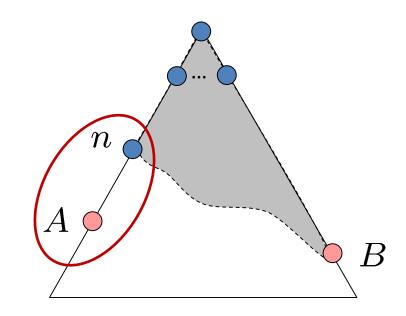


## Optimality of A\* Tree Search: Blocking



#### Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
  - 1. f(n) is less or equal to f(A)



$$f(n) = g(n) + h(n)$$

$$f(n) \le g(A)$$

$$g(A) = f(A)$$

Definition of f-cost Admissibility of h h = 0 at a goal


## Optimality of A\* Tree Search: Blocking



35

#### Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
  - 1. f(n) is less or equal to f(A)
  - 2. f(A) is less than f(B)



$$g(A) < g(B)$$

B is suboptimal

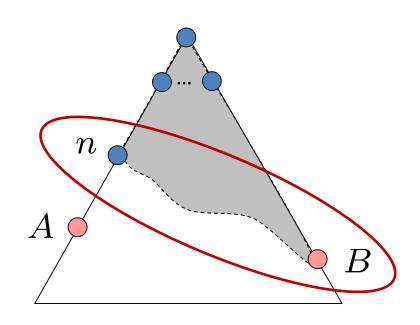
h = 0 at a goal

## Optimality of A\* Tree Search: Blocking



#### Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
  - 1. f(n) is less or equal to f(A)
  - 2. f(A) is less than f(B)
  - 3. *n* expands before B
- All ancestors of A expand before B
- A expands before B
- A\* search is optimal



$$f(n) \le f(A) < f(B)$$

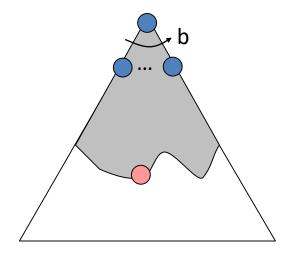


# Properties of A\*

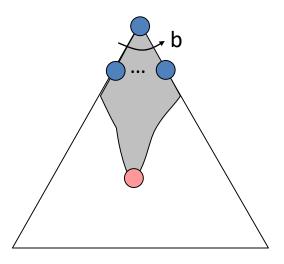
# Properties of A\*



#### **Uniform-Cost**



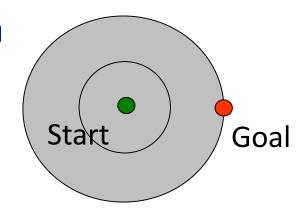




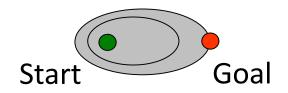
#### UCS vs A\* Contours



 Uniform-cost expands equally in all "directions"



 A\* expands mainly toward the goal, but does hedge its bets to ensure optimality



# Video of Demo Contours (Empty) -- UCS





# Video of Demo Contours (Empty) -- Greedy





# Video of Demo Contours (Empty) – A\*





# Video of Demo Contours (Pacman Small Maze) – A\*





# Comparison





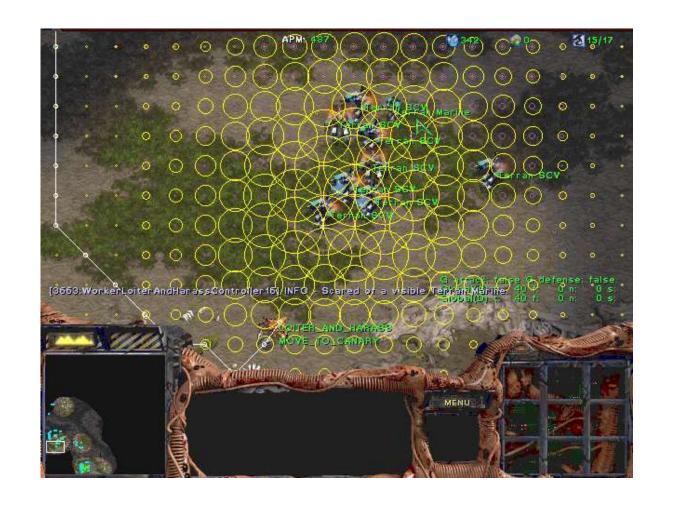
Greedy

**Uniform Cost** 

**A**\*

# A\* Applications





# A\* Applications



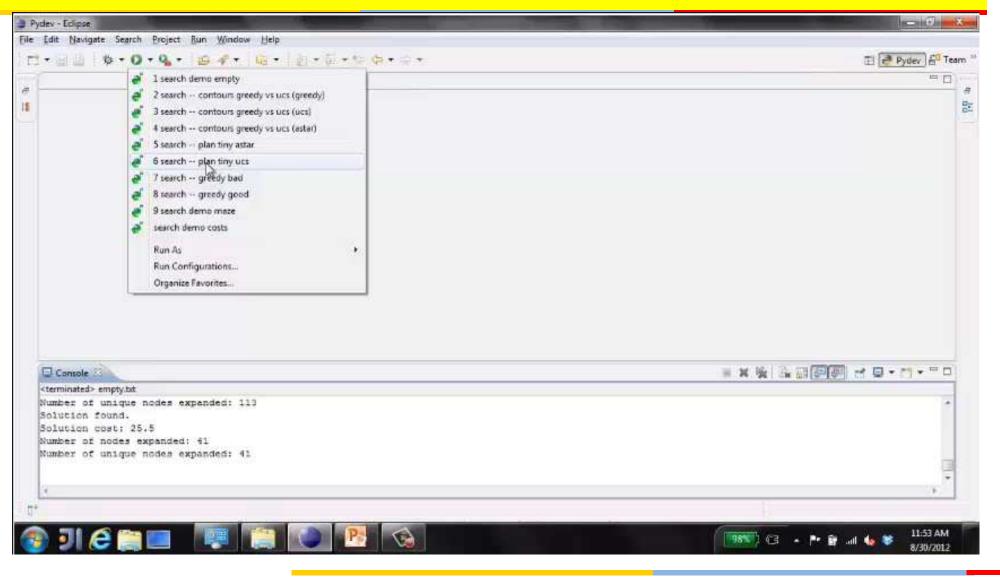
- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis

•



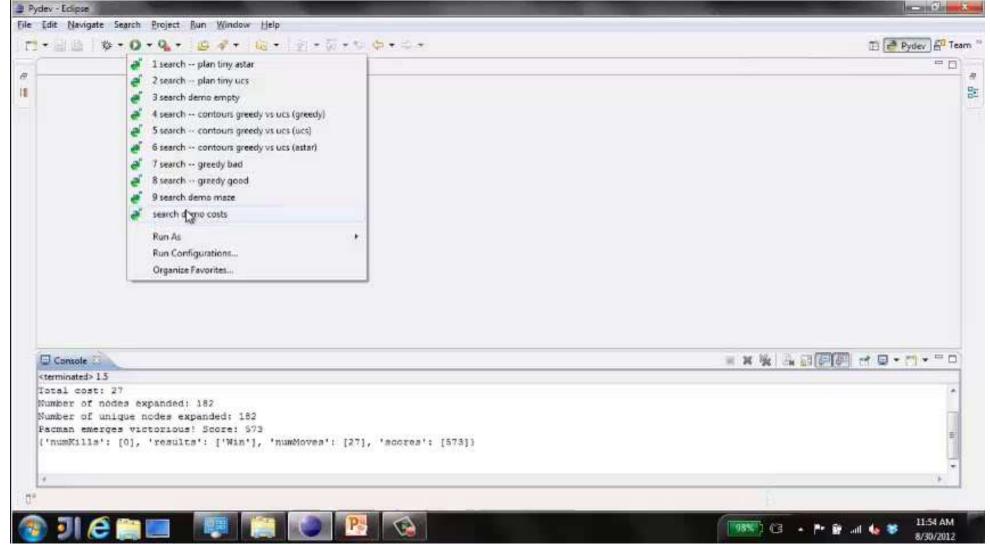
# Video of Demo Pacman (Tiny Maze) – UCS / A\*





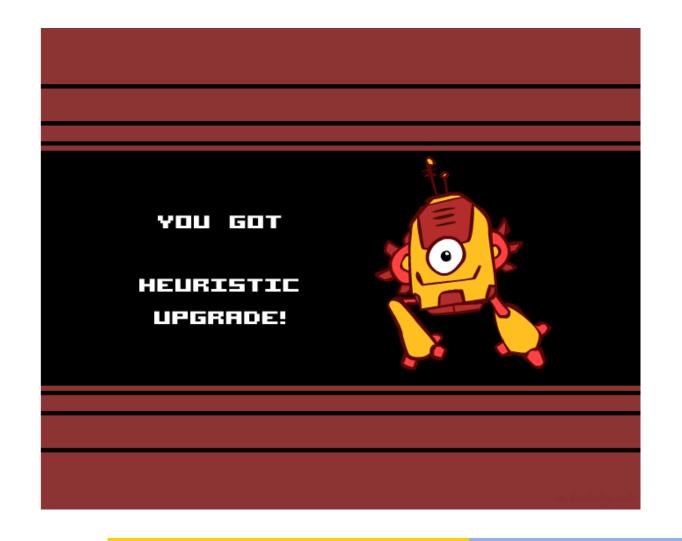
# Video of Demo Empty Water Shallow/Deep – **Guess Algorithm**





# **Creating Heuristics**

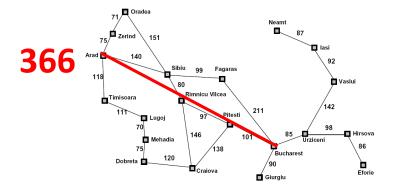





# Creating Admissible Heuristics



- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to relaxed problems, where new actions are available

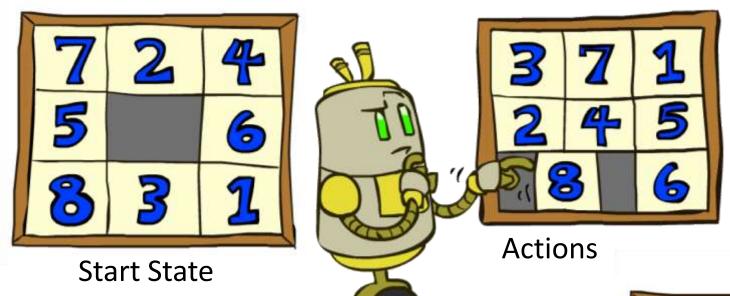




Inadmissible heuristics are often useful too

# Example: 8 Puzzle






• What are the states?

- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

**Goal State** 

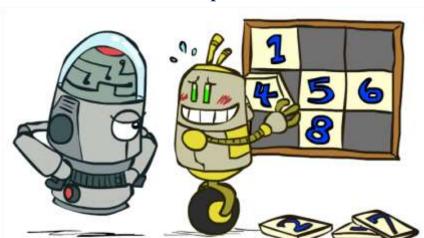


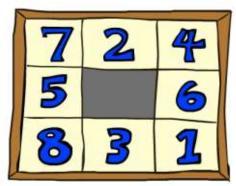
#### 8 Puzzle I

BITS Pilani

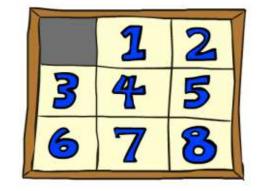
Pilani Campus

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- h(start) = 8
- This is a *relaxed-problem* heuristic









**Goal State** 

Average nodes expanded				
when the optimal path has				
4 steps	8 steps	12 steps		
112	6,300	$3.6 \times 10^6$		
13	39	227		

UCS

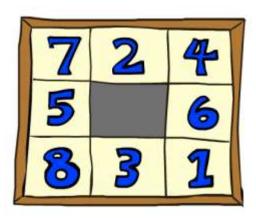
**TILES** 

#### 8 Puzzle II




53

- What if we had an easier 8puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total Manhattan distance
- Why is it admissible?
- h(start) = 3 + 1 + 2 + ... = 18







**Goal State** 


	when the optimal path has		
	4 steps	8 steps	12 steps
TILES	13	39	227
MANHATTAN	12	25	73

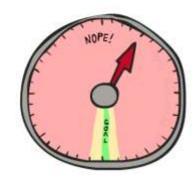
#### 8 Puzzle III



- How about using the *actual cost* as a heuristic?
  - Would it be admissible?
  - Would we save on nodes expanded?
  - What's wrong with it?







- With A\*: a trade-off between quality of estimate and work per node
  - As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself



# Semi-Lattice of Heuristics

### Trivial Heuristics, Dominance




• Dominance:  $h_a \ge h_c$  if

$$\forall n: h_a(n) \geq h_c(n)$$

- Heuristics form a semi-lattice:
  - Max of admissible heuristics is admissible

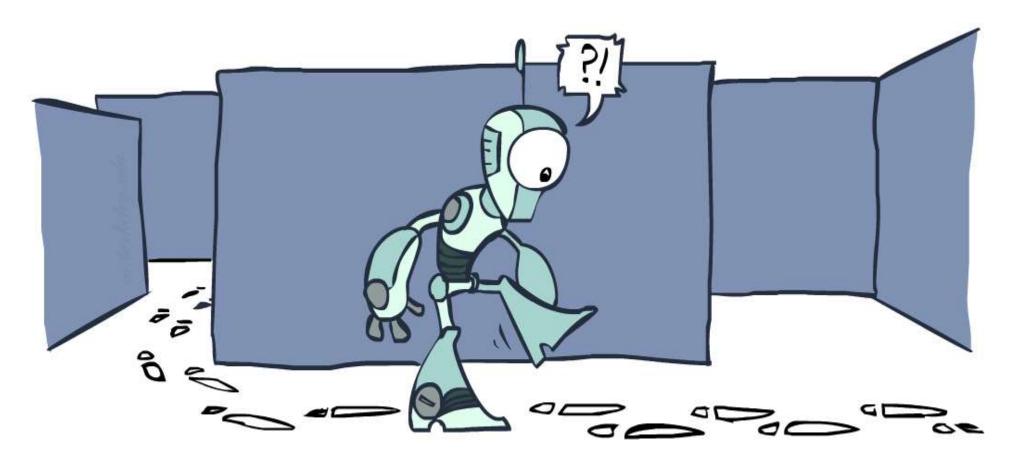
$$h(n) = \max(h_a(n), h_b(n))$$

- Trivial heuristics
  - Bottom of lattice is the zero heuristic (what does this give us?)
  - Top of lattice is the exact heuristic



# Graph Search







#### Tree Search: Extra Work!



• Failure to detect repeated states can cause exponentially more work.






## **Graph Search**



• In BFS, for example, we shouldn't bother expanding the circled nodes (why?)



# Graph Search




60

- Idea: never expand a state twice
- How to implement:
  - Tree search + set of expanded states ("closed set")
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set
- Important: store the closed set as a set, not a list
- Can graph search wreck completeness? Why/why not?
- How about optimality?


# A\* Graph Search Gone Wrong?



#### State space graph




#### Search tree



## **Consistency of Heuristics**



62



- Main idea: estimated heuristic costs ≤ actual costs
  - Admissibility: heuristic cost ≤ actual cost to goal
     h(A) ≤ actual cost from A to G
  - Consistency: heuristic "arc" cost ≤ actual cost for each arc
     h(A) h(C) ≤ cost(A to C)
- Consequences of consistency:
  - The f value along a path never decreases
     h(A) ≤ cost(A to C) + h(C)
  - A\* graph search is optimal

# Optimality of A\* Graph Search





# Optimality of A\* Graph Search




 Sketch: consider what A\* does with a consistent heuristic:

 Fact 1: In tree search, A\* expands nodes in increasing total f value (f-contours)

 Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally

Result: A\* graph search is optimal



# **Optimality**



- Tree search:
  - A\* is optimal if heuristic is admissible
  - UCS is a special case (h = 0)
- Graph search:
  - A\* optimal if heuristic is consistent
  - UCS optimal (h = 0 is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems



# A\*: Summary





# A\*: Summary



- A\* uses both backward costs and (estimates of) forward costs
- A\* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



#### Tree Search Pseudo-Code



```
function Tree-Search(problem, fringe) return a solution, or failure

fringe ← Insert(make-node(initial-state[problem]), fringe)

loop do

if fringe is empty then return failure

node ← remove-front(fringe)

if goal-test(problem, state[node]) then return node

for child-node in expand(state[node], problem) do

fringe ← insert(child-node, fringe)

end

end
```

## Graph Search Pseudo-Code



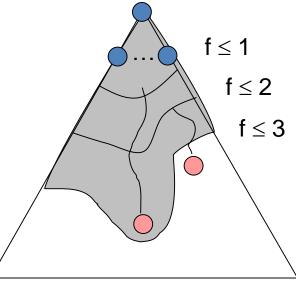
```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
   closed \leftarrow an empty set
   fringe \leftarrow Insert(Make-Node(Initial-state[problem]), fringe)
   loop do
       if fringe is empty then return failure
       node \leftarrow \text{REMOVE-FRONT}(fringe)
       if GOAL-TEST(problem, STATE[node]) then return node
       if STATE[node] is not in closed then
          add STATE[node] to closed
          for child-node in EXPAND(STATE[node], problem) do
              fringe \leftarrow INSERT(child-node, fringe)
          end
   end
```

# Optimality of A\* Graph Search



- Consider what A\* does:
  - Expands nodes in increasing total f value (f-contours) Reminder: f(n) = g(n) + h(n) = cost to n + heuristic
  - Proof idea: the optimal goal(s) have the lowest f value, so it must get expanded first

There's a problem with this argument. What are we assuming is true?

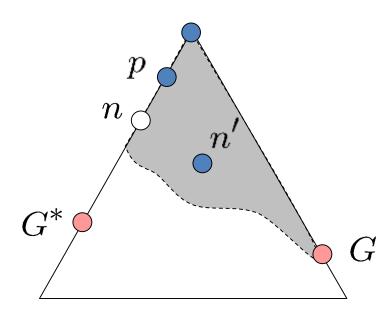


# Optimality of A\* Graph Search



#### Proof:

- New possible problem: some *n* on path to G\* isn't in queue when we need it, because some worse *n*' for the same state dequeued and expanded first (disaster!)
- Take the highest such *n* in tree
- Let *p* be the ancestor of *n* that was on the queue when *n*' was popped
- f(p) < f(n) because of consistency
- f(n) < f(n') because n' is suboptimal
- p would have been expanded before n'
- Contradiction!



#### Thank You





# Any more Queries